

Name: Aditya Mishra
Class:T.E.I.T

Roll No:34
Batch: A

In []: Experiment NO:**10**

In []: Aim:Business intelligencc mini project(Movie Rating Prediction)
LO:L01-L06

Theory:

Movie rating prediction is the task of predicting how a user will rate a movie based on their historical ratings and the characteristics of the movie. It is a classic example of a recommendation system, which aims to suggest relevant items to users based on their preferences and behavior.

One common approach to movie rating prediction is collaborative filtering, which involves using the ratings of other users who have similar preferences to predict the rating of a target user for a particular movie. Collaborative filtering can be done using either user-based or item-based approaches.

In the user-based approach, the similarity between users is computed based on their rating patterns, and the ratings of similar users are used to predict the rating of the target user for a particular movie. In the item-based approach, the similarity between movies is computed based on the rating patterns of users who have rated both movies, and the ratings of similar movies are used to predict the rating of the target user for a particular movie.

Another approach to movie rating prediction is content-based filtering, which involves using the characteristics of the movie to predict the rating of a target user. In this approach, the movie is described by a set of features, such as the genre, actors, director, and plot, and the rating of the target user is predicted based on the similarity between the features of the movie and the historical preferences of the user.

Machine learning techniques such as regression, decision trees, and neural networks can also be used for movie rating prediction. These techniques involve training a model on historical ratings and movie characteristics, and using the model to predict the rating of a target user for a particular movie.

Movie rating prediction is a challenging task due to the complexity of human preferences and the diversity of movies. However, it has important practical applications in movie recommendation systems and personalized advertising.

```
In [ ]: # Importing all the required libraries
```

```
In [1]: #Import the Libraries
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns
plt.style.use('seaborn-whitegrid')

from sklearn.impute import SimpleImputer

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error
from sklearn.preprocessing import LabelEncoder
```

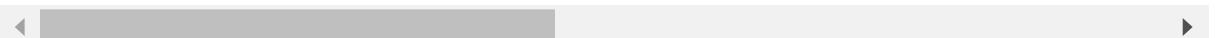
Loading the dataset

In [2]: #Load the Dataset

```
df_movies = pd.read_csv("kaggle_movie_dataset.csv", index_col="ID")
df_movies.head(10)
```

Out[2]:

	Unnamed: 0	Title	Year	IMDb	Netflix	Hulu	Prime Video	Disney+	Type	Directors
ID										
1	0	Inception	2010	8.8	1	0	0	0	0	Christopher Nolan
2	1	The Matrix	1999	8.7	1	0	0	0	0	Lana Wachowski,Lilly Wachowski
3	2	Avengers: Infinity War	2018	8.5	1	0	0	0	0	Anthony Russo,Joe Russo
4	3	Back to the Future	1985	8.5	1	0	0	0	0	Robert Zemeckis
5	4	The Good, the Bad and the Ugly	1966	8.8	1	0	1	0	0	Sergio Leone
6	5	Spider-Man: Into the Spider-Verse	2018	8.4	1	0	0	0	0	Bob Persichetti,Peter Ramsey,Rodney Rothman
7	6	The Pianist	2002	8.5	1	0	1	0	0	Roman Polanski
8	7	Django Unchained	2012	8.4	1	0	0	0	0	Quentin Tarantino
9	8	Raiders of the Lost Ark	1981	8.4	1	0	0	0	0	Steven Spielberg
10	9	Inglourious Basterds	2009	8.3	1	0	0	0	0	Quentin Tarantino



In [3]: df_movies.shape

Out[3]: (16744, 14)

Cleaning the dataset

In [4]: `#Remove "Unnamed"
df_movies=df_movies.drop("Unnamed: 0",axis=1)
df_movies.head()`

Out[4]:

ID	Title	Year	IMDb	Netflix	Hulu	Prime Video	Disney+	Type	Directors	Genres
1	Inception	2010	8.8	1	0	0	0	0	Christopher Nolan	Action, Adventure, Sci-Fi, Thriller
2	The Matrix	1999	8.7	1	0	0	0	0	Lana Wachowski, Lilly Wachowski	Action, Sci-Fi, Thriller
3	Avengers: Infinity War	2018	8.5	1	0	0	0	0	Anthony Russo, Joe Russo	Action, Adventure, Sci-Fi
4	Back to the Future	1985	8.5	1	0	0	0	0	Robert Zemeckis	Adventure, Comedy
5	The Good, the Bad and the Ugly	1966	8.8	1	0	1	0	0	Sergio Leone	Western

In [5]: `#Description overview of the data
df_movies.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 16744 entries, 1 to 16744
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Title            16744 non-null   object 
 1   Year             16744 non-null   int64  
 2   IMDb             16173 non-null   float64
 3   Netflix          16744 non-null   int64  
 4   Hulu             16744 non-null   int64  
 5   Prime Video      16744 non-null   int64  
 6   Disney+          16744 non-null   int64  
 7   Type             16744 non-null   int64  
 8   Directors         16018 non-null   object 
 9   Genres           16469 non-null   object 
 10  Country          16309 non-null   object 
 11  Language         16145 non-null   object 
 12  Runtime           16152 non-null   float64
dtypes: float64(2), int64(6), object(5)
memory usage: 1.8+ MB
```

finding total number of null values

In [6]: #Check for total no.of null values in each column
df_movies.isna().sum()

Out[6]:

Title	0
Year	0
IMDb	571
Netflix	0
Hulu	0
Prime Video	0
Disney+	0
Type	0
Directors	726
Genres	275
Country	435
Language	599
Runtime	592
dtype:	int64

In [7]: #Removing the Target value
df_movies = df_movies[df_movies['IMDb'].notna()]

```
In [8]: pip install missingno
```

```
Requirement already satisfied: missingno in c:\users\lenovo\anaconda3\lib\site-packages (0.5.2)
Note: you may need to restart the kernel to use updated packages.
```

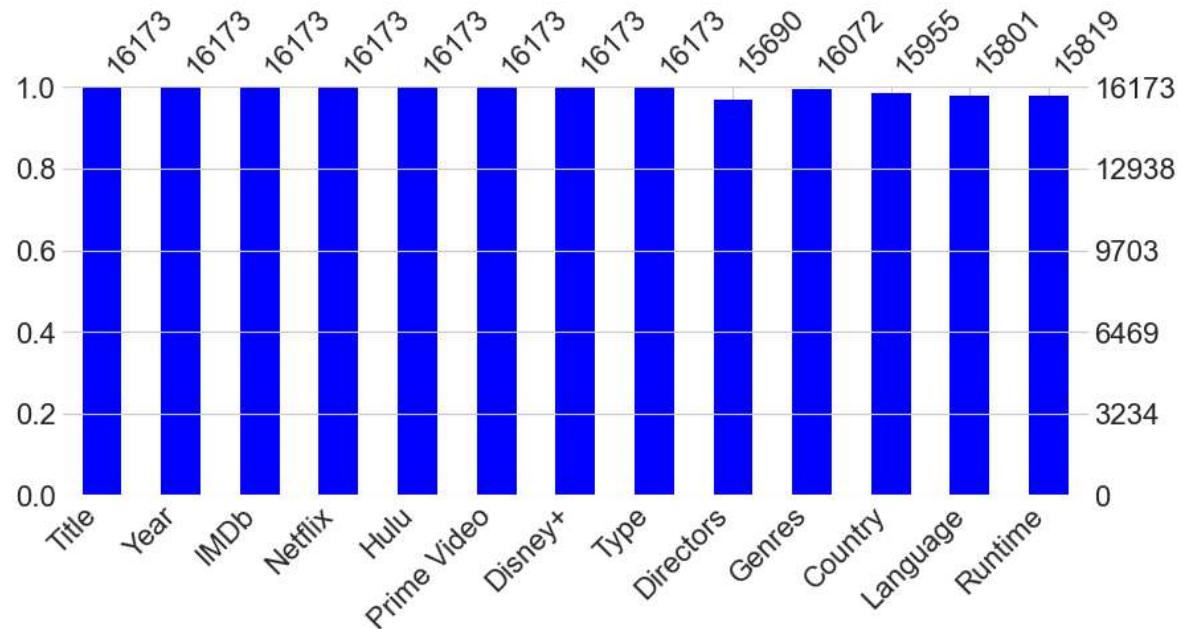
```
Requirement already satisfied: matplotlib in c:\users\lenovo\anaconda3\lib\site-packages (from missingno) (3.5.2)
Requirement already satisfied: seaborn in c:\users\lenovo\anaconda3\lib\site-packages (from missingno) (0.11.2)
Requirement already satisfied: numpy in c:\users\lenovo\anaconda3\lib\site-packages (from missingno) (1.21.5)
Requirement already satisfied: scipy in c:\users\lenovo\anaconda3\lib\site-packages (from missingno) (1.9.1)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib->missingno) (2.8.2)
Requirement already satisfied: cycler>=0.10 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib->missingno) (0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib->missingno) (1.4.2)
Requirement already satisfied: pillow>=6.2.0 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib->missingno) (9.2.0)
Requirement already satisfied: pyparsing>=2.2.1 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib->missingno) (3.0.9)
Requirement already satisfied: packaging>=20.0 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib->missingno) (21.3)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\lenovo\anaconda3\lib\site-packages (from matplotlib->missingno) (4.25.0)
Requirement already satisfied: pandas>=0.23 in c:\users\lenovo\anaconda3\lib\site-packages (from seaborn->missingno) (1.4.4)
Requirement already satisfied: pytz>=2020.1 in c:\users\lenovo\anaconda3\lib\site-packages (from pandas>=0.23->seaborn->missingno) (2022.1)
Requirement already satisfied: six>=1.5 in c:\users\lenovo\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib->missingno) (1.16.0)
```

```
In [9]: import missingno as msno
```

Representation of missing data

In [10]: #Visualizing the amount of missing data
`msno.bar(df_movies ,color='blue', figsize=(10, 4))`

Out[10]: <AxesSubplot:>



Removal of null values

In [11]: #Dropping all the rows(entries) where there are celss with no data
`df_movies.dropna(axis=0, how='any', inplace=True)`

`df_movies.isna().sum()`

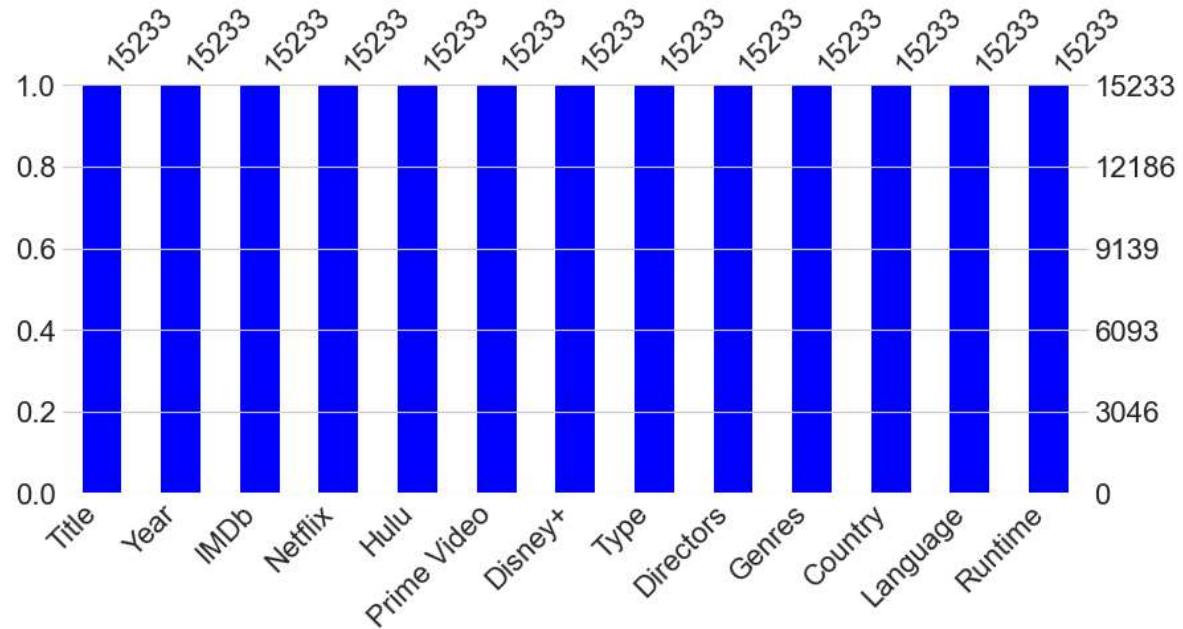
Out[11]:

Feature	Count
Title	0
Year	0
IMDb	0
Netflix	0
Hulu	0
Prime Video	0
Disney+	0
Type	0
Directors	0
Genres	0
Country	0
Language	0
Runtime	0

`dtype: int64`

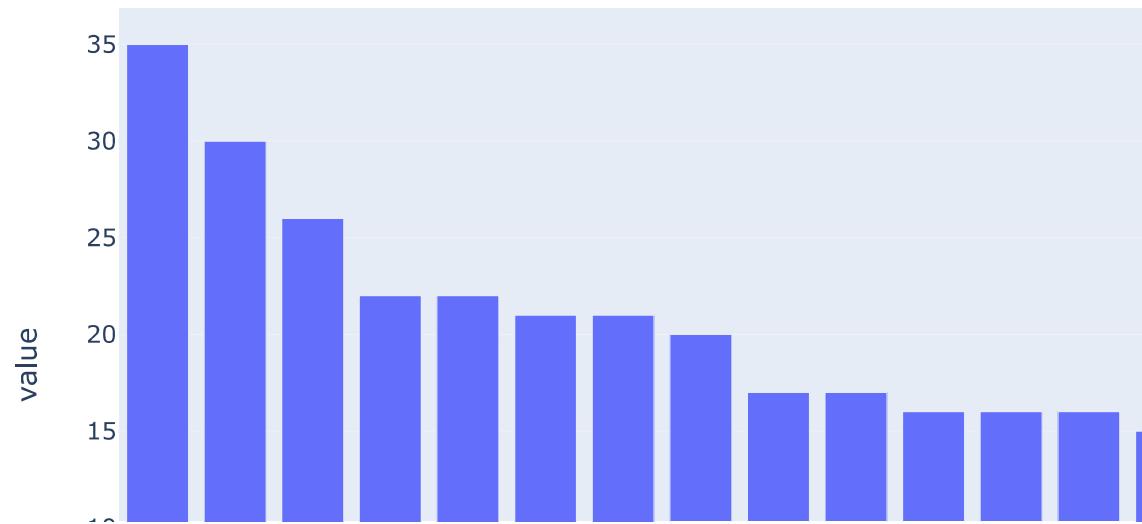
In [12]: #Visualizing if there is anymore missing data or not
msno.bar(df_movies ,color='blue', figsize=(10, 4))

Out[12]: <AxesSubplot:>



Bar Graph Representation of Director versus number of movies made by director

```
In [13]: import plotly.express as px
director = df_movies.dropna(axis=0, subset=['Directors'])
dfg = director['Directors'].value_counts().head(20)
fig = px.bar(dfg)
fig.show()
```



Representation of top movie of Each year

```
In [14]: import plotly.figure_factory as ff
import plotly.offline as py
rating = df_movies.dropna(axis=0, subset=['IMDb'])
rating = rating.sort_values('IMDb', ascending=False).drop_duplicates(['Year'])
data = ff.create_table(rating)
py.iplot(data)
```

Title	Year	IMDb	Netflix	Hulu	Prime	Vudu	Disney+	Type	Director	Genres
Down, But Not Out	2015	9.3	0	0	1	0	0	0	Miguel Góñlez	Drama
Steven Ballmer's Home Entertainment Center	1989	9.3	0	0	1	0	0	0	Tom McLoughlin	Documentary
Square One	2019	9.3	0	0	1	0	0	0	Danny Wu	Documentary
Bounty	2011	9.3	0	0	1	0	0	0	Roger Donaldson	Action
Natsamra	2016	9.1	1	0	0	0	0	0	Mahesh Manjrekar	Fiction
Finding Fa	2013	9.1	0	0	1	0	0	0	Chris Leslie	Drama
Where's Daddy?	2007	9.1	0	0	1	0	0	0	Rel Dowd	Documentary
The Dark	2008	9.0	0	1	0	0	0	0	Christopher Nolan	Action
8 Wheels	2006	8.9	0	0	1	0	0	0	Tyrone D. Wilson	Documentary
The Creat	2012	8.9	0	0	1	0	0	0	Laura Geller	Documentary

Scatter representation of IMDb rating and movie

```
In [15]: import matplotlib.pyplot as plt
import mplcursors

# Get the movie names and IMDb ratings for the first 50 movies
movie_names = df_movies['Title'][:50]
imdb_ratings = df_movies['IMDb'][:50]

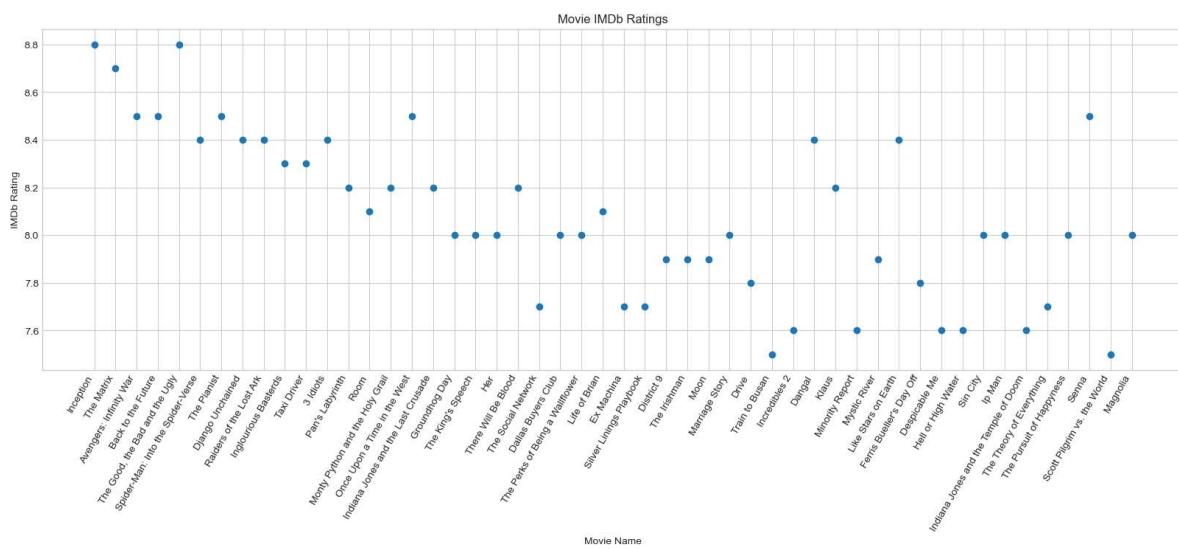
# Set the figure size and create a scatter plot with the movie names and IMDb
fig, ax = plt.subplots(figsize=(20, 6))
scatter = ax.scatter(movie_names, imdb_ratings)

# Rotate the x-axis labels for better readability
plt.xticks(rotation=60, ha='right')

# Add axis labels and a title
ax.set_xlabel('Movie Name')
ax.set_ylabel('IMDb Rating')
ax.set_title('Movie IMDb Ratings')

# Add a scroll bar to the x-axis
mplcursors.cursor(scatter, hover=True)

# Show the plot
plt.show()
```



In [16]: `gen = df_movies['Genres'].value_counts()
gen`

Out[16]:

Drama	1225
Documentary	1034
Comedy	905
Comedy, Drama	428
Horror	420
...	
Animation, Adventure, Drama, Family, Fantasy, Musical, Mystery, Romance	1
Adventure, Biography, Drama	1
Documentary, Biography, Family, History, Sport	1
Musical, Comedy, Romance	1
Comedy, Family, Adventure, Fantasy, Sci-Fi	1

Name: Genres, Length: 1824, dtype: int64

In [17]: `df_movies1 = pd.DataFrame(list(gen.items()))
df_movies1 = df_movies1.rename(columns={0: 'Genres', 1: 'Count'})
df_movies1`

Out[17]:

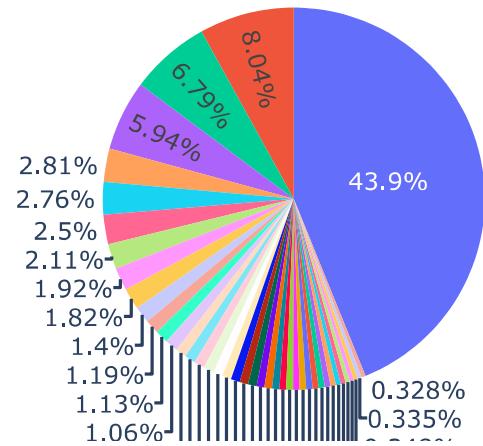
	Genres	Count
0	Drama	1225
1	Documentary	1034
2	Comedy	905
3	Comedy, Drama	428
4	Horror	420
...
1819	Animation, Adventure, Drama, Family, Fantasy, Music...	1
1820	Adventure, Biography, Drama	1
1821	Documentary, Biography, Family, History, Sport	1
1822	Musical, Comedy, Romance	1
1823	Comedy, Family, Adventure, Fantasy, Sci-Fi	1

1824 rows × 2 columns

In [18]: `#Pie chart representation of Genres`

```
In [19]: import plotly.express as px  
df_movies1.loc[df_movies1['Count'] < 50, 'Genres'] = 'Other' # Represent only  
fig = px.pie(df_movies1, values='Count', names='Genres', title='Genre Distribu  
fig.show()
```

Genre Distribution



```
In [20]: # Assuming your data is stored in a pandas DataFrame called "df"  
  
# Create a new column to store the director's average rating  
df_movies['director_avg_rating'] = 0.0  
  
# Loop through each director in the data  
for director in df_movies['Directors'].unique():  
    # Calculate the average rating of the director's movies  
    avg_rating = df_movies.loc[df_movies['Directors'] == director, 'IMDb'].mean()  
    # Update the director_avg_rating column with the average rating  
    df_movies.loc[df_movies['Directors'] == director, 'director_avg_rating'] =
```

In [21]: df_movies.head()

Out[21]:

ID	Title	Year	IMDb	Netflix	Hulu	Prime Video	Disney+	Type	Directors	Genres
1	Inception	2010	8.8	1	0	0	0	0	Christopher Nolan	Action, Adventure, Sci-Fi, Thriller
2	The Matrix	1999	8.7	1	0	0	0	0	Lana Wachowski, Lilly Wachowski	Action, Sci-Fi, Thriller
3	Avengers: Infinity War	2018	8.5	1	0	0	0	0	Anthony Russo, Joe Russo	Action, Adventure, Sci-Fi
4	Back to the Future	1985	8.5	1	0	0	0	0	Robert Zemeckis	Adventure, Comedy
5	The Good, the Bad and the Ugly	1966	8.8	1	0	1	0	0	Sergio Leone	Western

Choosing Target Variable and features

Target variable --> IMDb

**Categorial variable-->
Year, Runtime, Directors, Genres, Country, Language**

```
In [28]: # Choose target and features
y = df_movies.IMDb

X = df_movies.drop(['IMDb'], axis=1)

#Split the data for train and test
X_train_full, X_test_full, y_train, y_test = train_test_split(X, y, random_state=42, test_size=0.2)

#List of Categorical columns to be used as features
cat_cols=["Directors","Genres","Country","Language",'director_avg_rating']

#List of Numerical columns to be used as features
numerical_cols = ['Year','Runtime']

#Keep selected columns only
my_cols = numerical_cols + cat_cols
X_train = X_train_full[my_cols].copy()
X_test = X_test_full[my_cols].copy()
from sklearn.preprocessing import LabelEncoder

# Convert categorical columns to numeric using LabelEncoder
le = LabelEncoder()
df_movies['Directors'] = le.fit_transform(df_movies['Directors'])

# Convert categorical columns to numeric using LabelEncoder
le = LabelEncoder()
df_movies['Genres'] = le.fit_transform(df_movies['Genres'])

# Convert categorical columns to numeric using LabelEncoder
le = LabelEncoder()
df_movies['Country'] = le.fit_transform(df_movies['Country'])

# Convert categorical columns to numeric using LabelEncoder
le = LabelEncoder()
df_movies['Language'] = le.fit_transform(df_movies['Language'])
```

In [29]: X_train.head()

Out[29]:

	Year	Runtime	Directors	Genres	Country	Language	director_avg_rating
--	------	---------	-----------	--------	---------	----------	---------------------

ID							
13576	1972	88.0	10337	224	1024	97	5.600000
9829	2011	93.0	7359	1115	1024	97	4.600000
1256	2020	104.0	10606	1580	1024	182	7.233333
9800	2017	149.0	1226	389	463	1032	5.100000
11212	1941	66.0	3935	173	1024	97	6.311111

In [30]: `X_test.head()`

Out[30]:

ID	Year	Runtime	Directors	Genres	Country	Language	director_avg_rating
8874	2016	90.0	8806	1807	1024	97	5.8
12421	2016	87.0	6848	915	1024	97	6.1
8191	2018	79.0	9795	1735	1024	165	4.6
12796	2017	120.0	1829	95	1024	97	6.4
4721	2018	136.0	10760	1172	185	64	7.0

In [31]: `y_train.head()`

Out[31]: ID

13576 5.6
9829 4.6
1256 6.4
9800 5.1
11212 7.0
Name: IMDb, dtype: float64

In [32]: `y_test.head()`

Out[32]: ID

8874 5.8
12421 6.1
8191 4.6
12796 6.4
4721 7.0
Name: IMDb, dtype: float64

Linear Regression

```
In [33]: import pandas as pd
from sklearn.preprocessing import LabelEncoder

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

# Create Linear regression object
lr = LinearRegression()

# Fit the model using the training data
lr.fit(X_train, y_train)

# Make predictions using the test data
y_pred = lr.predict(X_test)

# Calculate the accuracy of the model
accuracy = lr.score(X_test, y_test)
linear_regression_acc=accuracy*100
print("Accuracy from Linear Regression:", accuracy*100)
```

Accuracy from Linear Regression: 87.19203843226072

r2_score calculation using linear regression

```
In [34]: from sklearn.metrics import r2_score

# Predict on the training and testing sets
y_train_pred = lr.predict(X_train)
y_test_pred = lr.predict(X_test)

# Calculate the R-squared value on the training and testing sets
train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)

# Print the R-squared values
print("Training R-squared value: {:.2f}".format(train_r2))
print("Testing R-squared value: {:.2f}".format(test_r2))
```

Training R-squared value: 0.85

Testing R-squared value: 0.87

RandomForestRegressor

```
In [35]: import pandas as pd
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split

# Create a Random Forest regression object
rf = RandomForestRegressor( random_state=42)

# Fit the model using the training data
rf.fit(X_train, y_train)

# Make predictions using the test data
y_pred = rf.predict(X_test)

# Calculate the accuracy of the model
accuracy = rf.score(X_test, y_test)
random_forest_acc=accuracy*100
print("Accuracy from RandomForestRegressor:", accuracy*100)
```

Accuracy from RandomForestRegressor: 83.99257811349999

r2_score calculation using RandomForestRegressor

```
In [36]: from sklearn.metrics import r2_score

# Predict on the training and testing sets
y_train_pred = rf.predict(X_train)
y_test_pred = rf.predict(X_test)

# Calculate the R-squared value on the training and testing sets
train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)

# Print the R-squared values
print("Training R-squared value: {:.2f}".format(train_r2))
print("Testing R-squared value: {:.2f}".format(test_r2))
```

Training R-squared value: 0.97

Testing R-squared value: 0.84

DecisionTreeRegressor

```
In [37]: from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import r2_score

# Create a decision tree regressor
dt = DecisionTreeRegressor()

# Fit the model to the training data
dt.fit(X_train, y_train)

# Predict on the test data
y_pred = dt.predict(X_test)

# Calculate the R-squared score (a measure of model accuracy)
accuracy = r2_score(y_test, y_pred)
decision_tree_acc=accuracy*100

# Print the accuracy score
print('Accuracy from DecisionTreeRegressor:', accuracy*100)
```

Accuracy from DecisionTreeRegressor: 68.28417073839233

r2_score calculation using DecisionTreeRegressor

```
In [38]: from sklearn.metrics import r2_score

# Predict on the training and testing sets
y_train_pred = dt.predict(X_train)
y_test_pred = dt.predict(X_test)

# Calculate the R-squared value on the training and testing sets
train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)

# Print the R-squared values
print("Training R-squared value: {:.2f}".format(train_r2))
print("Testing R-squared value: {:.2f}".format(test_r2))
```

Training R-squared value: 1.00
 Testing R-squared value: 0.68

```
In [39]: pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\lenovo\anaconda3\lib\site-packages (1.7.4)
Requirement already satisfied: numpy in c:\users\lenovo\anaconda3\lib\site-packages (from xgboost) (1.21.5)
Requirement already satisfied: scipy in c:\users\lenovo\anaconda3\lib\site-packages (from xgboost) (1.9.1)
Note: you may need to restart the kernel to use updated packages.
```

xgboost

```
In [40]: import xgboost as xgb
from sklearn.metrics import r2_score

# Create an XGBoost model with default parameters
xgb_model = xgb.XGBRegressor()

# Fit the model to the training data
xgb_model.fit(X_train, y_train)

# Predict on the test data
y_pred = xgb_model.predict(X_test)

# Calculate the R-squared score (a measure of model accuracy)
accuracy = r2_score(y_test, y_pred)
xgb_acc=accuracy*100

# Print the accuracy score
print('Accuracy:', accuracy*100)
```

Accuracy: 84.39630977667795

r2_score calculation using xgboost

```
In [41]: from sklearn.metrics import r2_score

# Predict on the training and testing sets
y_train_pred = xgb_model.predict(X_train)
y_test_pred = xgb_model.predict(X_test)

# Calculate the R-squared value on the training and testing sets
train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)

# Print the R-squared values
print("Training R-squared value: {:.2f}".format(train_r2))
print("Testing R-squared value: {:.2f}".format(test_r2))
```

Training R-squared value: 0.93
Testing R-squared value: 0.84

```
In [42]: print(y_pred)
```

```
[5.811812 6.016548 4.6789927 ... 4.188054 6.255905 7.219863 ]
```

Individual representation of each model

```
In [43]: # Predict on the training and testing sets
y_train_pred = lr.predict(X_train)
y_test_pred = lr.predict(X_test)

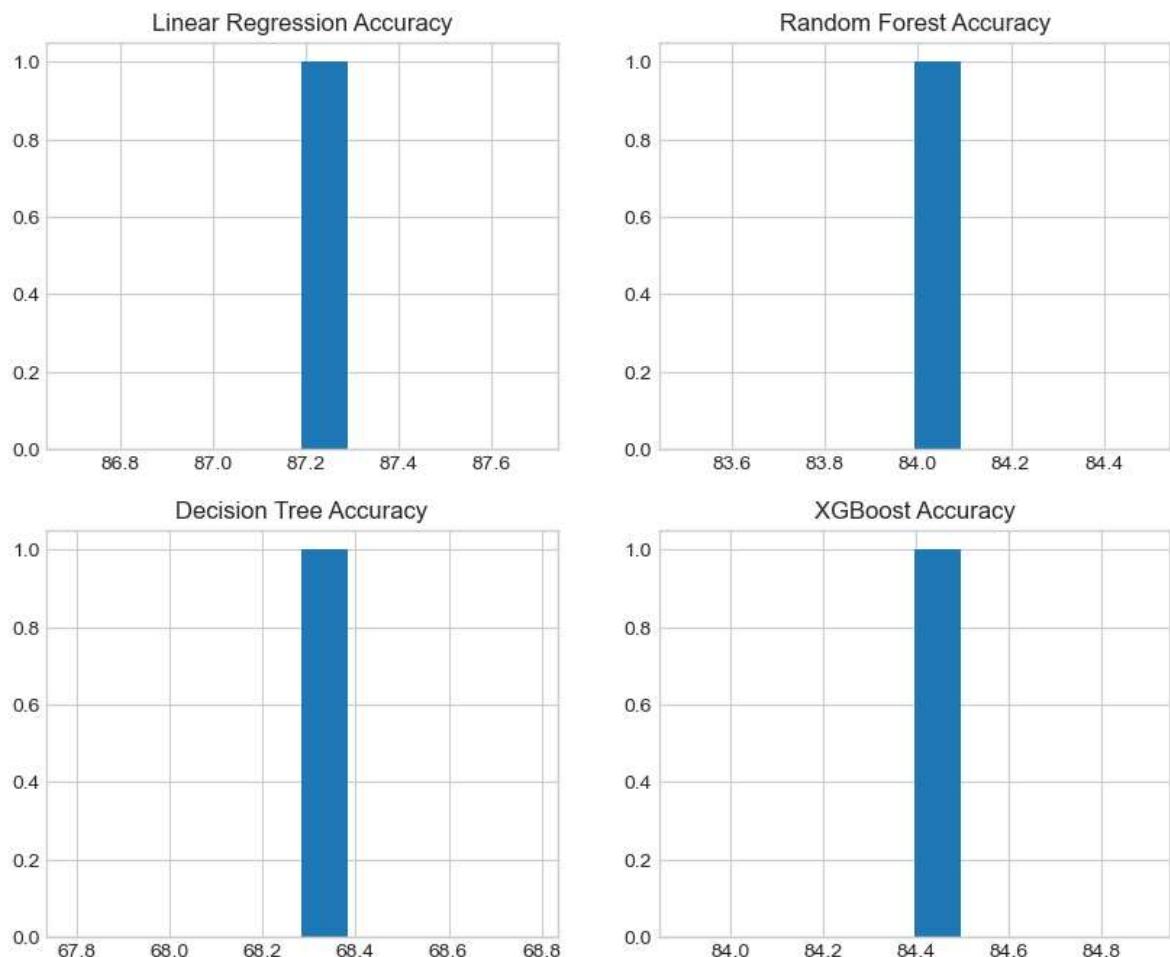
# Calculate the R-squared value on the training and testing sets
train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)

# Print the R-squared values
print("Training R-squared value: {:.2f}".format(train_r2))
print("Testing R-squared value: {:.2f}".format(test_r2))
```

Training R-squared value: 0.85
Testing R-squared value: 0.87

```
In [44]: # Create subplots for each model's histogram
fig, axs = plt.subplots(2, 2, figsize=(10, 8))
axs[0, 0].hist(linear_regression_acc, bins=10)
axs[0, 0].set_title('Linear Regression Accuracy')
axs[0, 1].hist(random_forest_acc, bins=10)
axs[0, 1].set_title('Random Forest Accuracy')
axs[1, 0].hist(decision_tree_acc, bins=10)
axs[1, 0].set_title('Decision Tree Accuracy')
axs[1, 1].hist(xgb_acc, bins=10)
axs[1, 1].set_title('XGBoost Accuracy')

# Show the plot
plt.show()
```



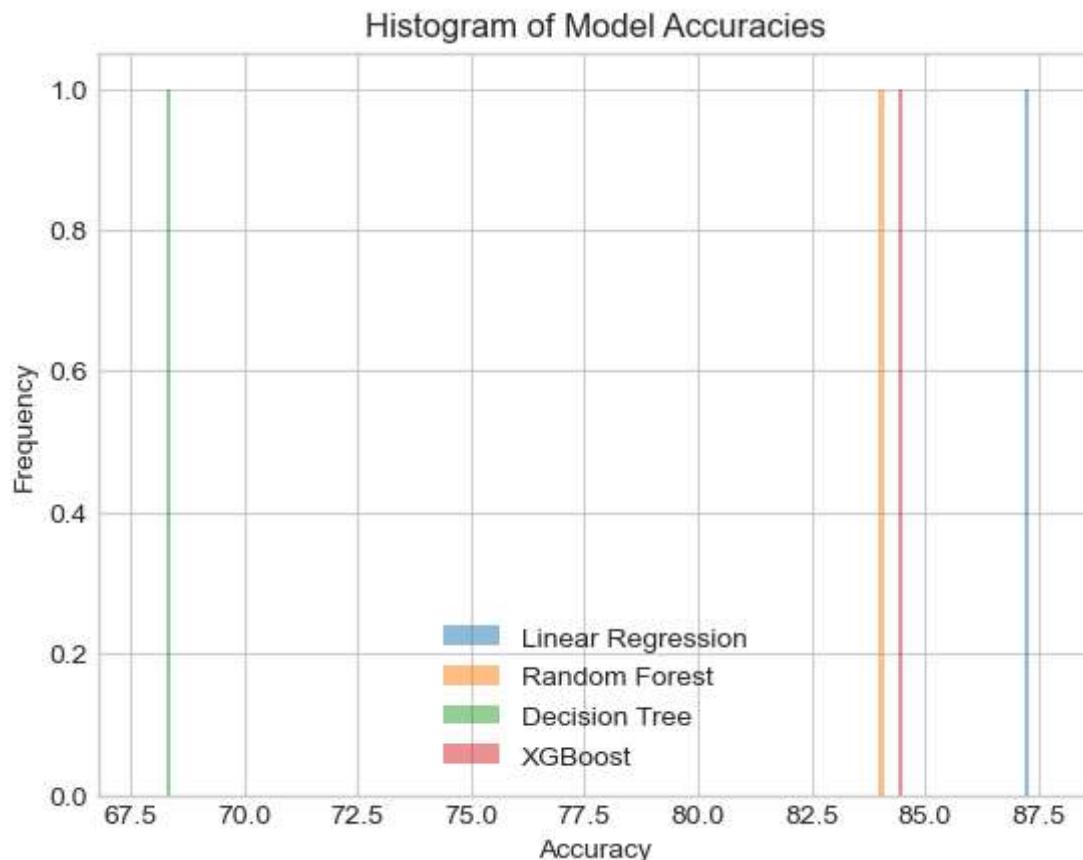
Subplot representation

In [45]:

```
# Create a histogram for each model
plt.hist(linear_regression_acc, bins=10, alpha=0.5, label='Linear Regression')
plt.hist(random_forest_acc, bins=10, alpha=0.5, label='Random Forest')
plt.hist(decision_tree_acc, bins=10, alpha=0.5, label='Decision Tree')
plt.hist(xgb_acc, bins=10, alpha=0.5, label='XGBoost')

# Add Labels and Legend
plt.xlabel('Accuracy')
plt.ylabel('Frequency')
plt.title('Histogram of Model Accuracies')
plt.legend()

# Show the plot
plt.show()
```



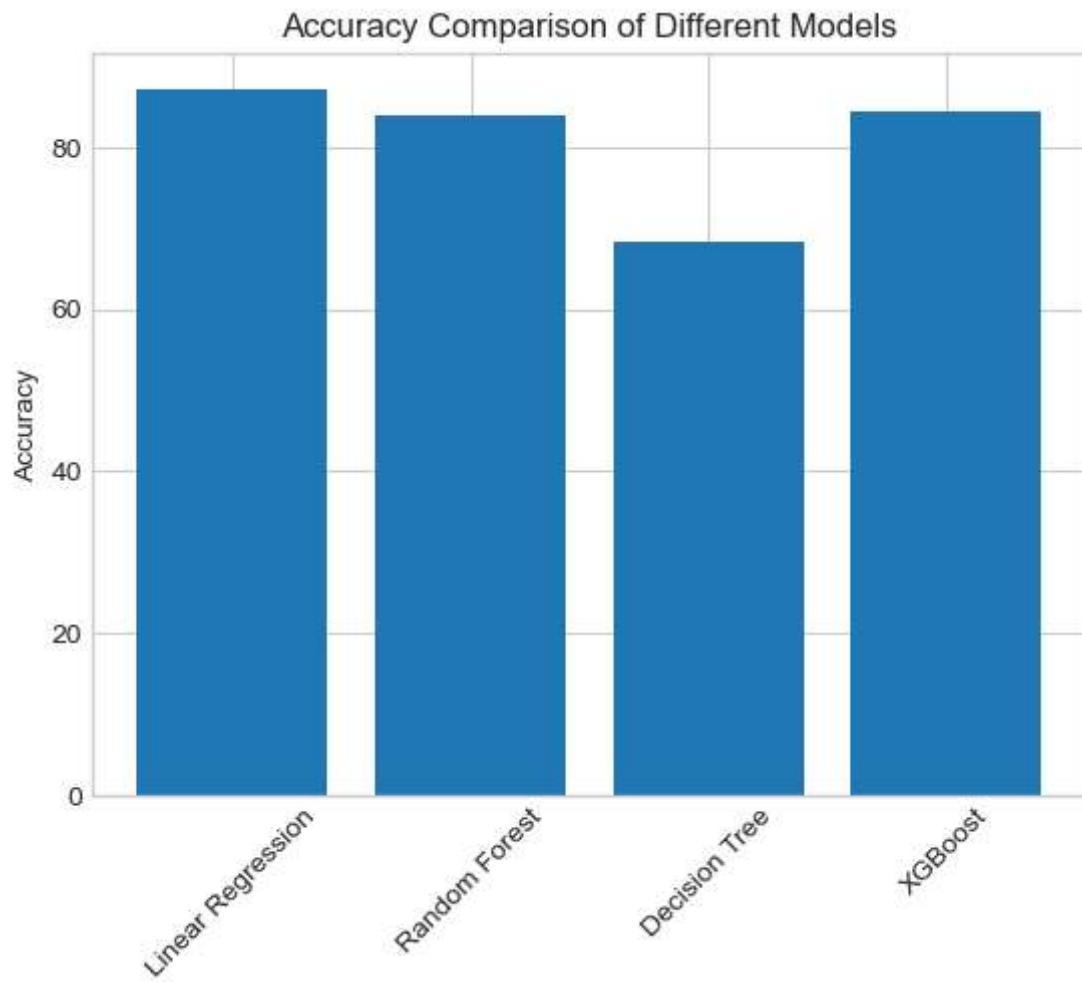
Histogram representation of model accuracy

```
In [46]: # Create a list of model names
models = ['Linear Regression', 'Random Forest', 'Decision Tree', 'XGBoost']
fig, ax = plt.subplots()
accuracy = [linear_regression_acc, random_forest_acc, decision_tree_acc, xgb_acc]

# Create a histogram with the sample data
ax.bar(models, accuracy)

# Add labels and title
ax.set_xticks(np.arange(len(models)))
ax.set_xticklabels(models, rotation=45)
ax.set_ylabel('Accuracy')
ax.set_title('Accuracy Comparison of Different Models')

# Show the plot
plt.show()
```



```
In [47]: import warnings
warnings.filterwarnings("ignore", message="X does not have valid feature names")
```

```
In [48]: import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LinearRegression

# Encode the categorical variables using Label encoding
le = LabelEncoder()
df_movies["Directors"] = le.fit_transform(df_movies["Directors"])
df_movies["Genres"] = le.fit_transform(df_movies["Genres"])
df_movies["Country"] = le.fit_transform(df_movies["Country"])
df_movies["Language"] = le.fit_transform(df_movies["Language"])

# Separate the features and the target variable
X = df_movies[["Year", "Runtime", "Directors", "Genres", "Country", "Language"]]
y = df_movies["IMDb"]

# Create Linear regression object
lr = LinearRegression()

# Fit the model using the entire dataset
lr.fit(X, y)

# Make a prediction using the input values
input_values = [[2020, 100, 205, 2, 401, 2, 8.2]]
predicted_rating = lr.predict(input_values)

# Print the predicted rating
print("Predicted IMDb rating:", predicted_rating[0])
```

Predicted IMDb rating: 8.13117171343653

Conclusion: From the above experiment i have learned how to implement miniproject on Movie Rating Prediction Project and i also learned how to use different machine learning model in this experiment i have achieved P01,P02,P03,P04,P05,P08,P09,P010,P012